

Exercise 2

Introduction to High-Performance Computing
WS 2019

Julian Miller

Tim Cramer

Sandra Wienke

contact@hpc.rwth-aachen.de

Recap: $O(N^3)/O(N^2)$ – Example

- Examples for $O(N^3)/O(N^2)$ algorithms are complex
- To simplify, we use an $O(N^2)/O(N)$ algorithm to illustrate possible optimization:

```
for(i = 0; i < N; ++i )  
{  
  for(j=0; j < N; ++j)  
  {  
    sum += a[i] * b[j];  
  }  
}
```


b is loaded from memory N times
a is loaded once per inner loop

Thus memory traffic amounts to $N(N+1)$ words.

- Loop unroll & jam will reduce this to $N(N/m + 1)$
→ with m-way unrolling


Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i)
{
    for(j=0; j < N; ++j)
    {
        sum += a[i] * b[j];
    }
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i)
{
    for(j=0; j < N; ++j)
    {
        sum += a[i] * b[j];
    }
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )  
{  
  for(j=0; j < N; ++j)  
  {  
    sum += a[i] * b[j];  
  }  
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )  
{  
  for(j=0; j < N; ++j)  
  {  
    sum += a[i] * b[j];  
  }  
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )  
{  
  for(j=0; j < N; ++j)  
  {  
    sum += a[i] * b[j];  
  }  
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )
{
  for(j=0; j < N; ++j)
  {
    sum += a[i] * b[j];
  }
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )
{
  for(j=0; j < N; ++j)
  {
    sum += a[i] * b[j];
  }
}
```




Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i )  
{  
  for(j=0; j < N; ++j)  
  {  
    sum += a[i] * b[j];  
  }  
}
```



Recap: $O(N^3)/O(N^2)$ – Example

```
for(i = 0; i < N; ++i)
{
    for(j=0; j < N; ++j)
    {
        sum += a[i] * b[j];
    }
}
```



Recap: $O(N^3)/O(N^2)$ – Example

■ Blocking the inner loop: —————→

→ With blocking factor B

■ Introduces two effects

→ Array b is only loaded once from memory now

→ as long as factor B is small enough that parts loaded from b fit into the cache and stay there as long as needed


→ Array a is loaded from memory N/B times instead of once

■ Effective memory traffic: $N(N/B+1)$

■ Blocking is the method of choice here, as the blocking factor can be increased to higher values than the unrolling factor


```
for(jj=0; jj < N; jj+=B)
{
    jstart=jj; jend=jj+B;
    for(i = 0; i < N; ++i )
    {
        for(j=jstart; j < jend; ++j)
        {
            sum += a[i] * b[j];
        }
    }
}
```

Recap: $O(N^3)/O(N^2)$ – Example




```
for(jj=0; jj < N; jj+=B)
{
    jstart=jj; jend=jj+B;
    for(i = 0; i < N; ++i )
    {
        for(j=jstart; j < jend; ++j)
        {
            sum += a[i] * b[j];
        }
    }
}
```

Recap: $O(N^3)/O(N^2)$ – Example




```
for(jj=0; jj < N; jj+=B)
{
  jstart=jj; jend=jj+B;
  for(i = 0; i < N; ++i )
  {
    for(j=jstart; j < jend; ++j)
    {
      sum += a[i] * b[j];
    }
  }
}
```

Recap: $O(N^3)/O(N^2)$ – Example




```
for(jj=0; jj < N; jj+=B)
{
  jstart=jj; jend=jj+B;
  for(i = 0; i < N; ++i )
  {
    for(j=jstart; j < jend; ++j)
    {
      sum += a[i] * b[j];
    }
  }
}
```

Recap: $O(N^3)/O(N^2)$ – Example




```
for(jj=0; jj < N; jj+=B)
{
  jstart=jj; jend=jj+B;
  for(i = 0; i < N; ++i )
  {
    for(j=jstart; j < jend; ++j)
    {
      sum += a[i] * b[j];
    }
  }
}
```


Recap: $O(N^3)/O(N^2)$ – Example



```
for(jj=0; jj < N; jj+=B)
{
  jstart=jj; jend=jj+B;
  for(i = 0; i < N; ++i )
  {
    for(j=jstart; j < jend; ++j)
    {
      sum += a[i] * b[j];
    }
  }
}
```

Recap: $O(N^3)/O(N^2)$ – Example



```
for(jj=0; jj < N; jj+=B)
{
  jstart=jj; jend=jj+B;
  for(i = 0; i < N; ++i )
  {
    for(j=jstart; j < jend; ++j)
    {
      sum += a[i] * b[j];
    }
  }
}
```

Exercise Tasks

1. **Norm calculation of a matrix**
2. Performance analysis tools
3. Balance Metric

1. Norm calculation of a matrix

Problem 1. Norm calculation of a matrix

Let $A = (a_{ij})_{i,j=1,\dots,n} \in \mathbb{R}^{n \times n}$ be a real matrix. The norms $\|\cdot\|_1$ and $\|\cdot\|_\infty$ are defined by:

$$\|A\|_1 := \max_{j=1,\dots,n} \sum_{i=1}^n |a_{ij}| \quad (\text{“maximum column sum”})$$

$$\|A\|_\infty := \max_{i=1,\dots,n} \sum_{j=1}^n |a_{ij}| \quad (\text{“maximum row sum”}).$$


This exercise may be done with pen and paper or using the C++ code template in `exercise2.tar.gz`

1. Norm calculation of a matrix

Preparation: Allocate memory for the matrix


- Idea: Store matrix as a 1D-block in memory

`double *B`



`double **A`

`double *B`



Same starting address: $B = *A$


Access elements by
 $A[x][y]$ or $B[x*n+y]$

1. Norm calculation of a matrix

Preparation: Allocate memory for the matrix


■ Example: $n = 4 \rightarrow 4 \times 4$ matrix

`double *B`



`double **A`

n pointers
to memory
location



`double *B`

Same starting address: $B = *A$

Access elements by
 $A[x][y]$ or $B[x*n+y]$

1. Norm calculation of a matrix

Preparation: Allocate memory for the matrix

■ Example: $n = 4 \rightarrow 4 \times 4$ matrix


`double *B`



16 elements

`double **A`

n pointers
to memory
location



`double *B`

Same starting address: $B = *A$

Access elements by
 $A[x][y]$ or $B[x*n+y]$


n

1. Norm calculation of a matrix


Preparation: Allocate memory for the matrix

■ Example: $n = 4 \rightarrow 4 \times 4$ matrix

`double *B`



`double **A`



`double *B`

Same starting address: $B = *A$


Access elements by
 $A[x][y]$ or $B[x*n+y]$

1. Norm calculation of a matrix

Preparation: Allocate memory for the matrix


■ Example: $n = 4 \rightarrow 4 \times 4$ matrix

`double *B`



`double **A`

n pointers
to memory
location



`double *B`

Same starting address: $B = *A$

Access elements by
 $A[x][y]$ or $B[x*n+y]$

Livedemo

1. Norm calculation of a matrix

■ Preparation: Allocate memory for the matrix

```
/* Allocate the complete matrix in one block */
```

```
double *B = new double[(long)n*n];
```

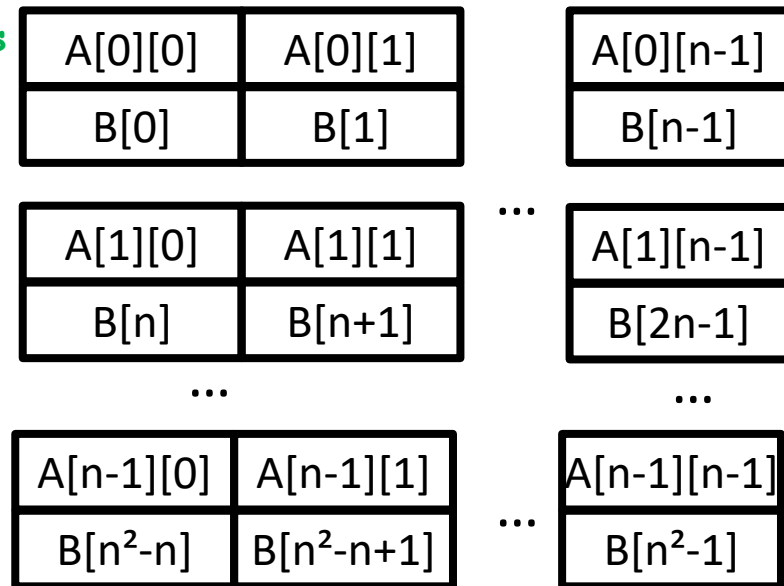


```
/* Allocate memory to store the address  
of the first element of each row */
```

```
double **A = new double*[n];
```

```
/* Place one pointer on each row  
of the matrix in the array A */
```

```
for (int i=0; i<n; ++i)  
    A[i] = &(B[(long)i*n]);
```



1. Norm calculation of a matrix

a) norm_max()

- a) Design an algorithm to compute the norm $\|A\|_{\infty}$ by passing through the memory consecutively.

Livedemo


1. Norm calculation of a matrix

a) norm_max()

```
double norm_max(double** const A, const int n) {
    double rowsum = 0., max_norm = -1.;
    /* TODO Implement the max norm using 2 for loops
       and the function double abs(double x) */
    for (int i = 0; i < n; ++i) {
        rowsum = 0.;
        for (int j = 0; j < n; ++j)
            rowsum += abs(A[i][j]);
        if (rowsum > max_norm)
            max_norm = rowsum;
    }
    return max_norm;
}
```


1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()




1. Norm calculation of a matrix

a) norm_max()



1. Norm calculation of a matrix

a) norm_max()



1. Norm calculation of a matrix

b) Calculate the time $T^{(1)}(n)$

b) Calculate the cache-miss-ratio ($\#cachemiss/\#memoryaccess$) and the time $T^{(1)}(n)$ of that algorithm in terms of T_A , T_M , T_C , l , and c where

T_A time for an arithmetic operation

T_M time for accessing the main memory

T_C time for accessing the cache

l size of a cache line (in elements)

c size of the cache (in elements).

Hint: Auxiliary variables are stored in registers. Assume that the computation of $|\cdot|$ and the comparison of two real numbers both take time T_A . Assume for all calculations that the matrix doesn't fit to the cache ($c \ll n \cdot n$) and the cache is initially cold (empty).

1. Norm calculation of a matrix

b) Calculate the time $T^{(1)}(n)$

- n^2 elements (the whole matrix) must be added together and every element must be examined by `abs()`: $2n^2 T_A$
- Comparing the row sums: nT_A
- Loading the matrix elements can be divided in two different actions:

→ Loading the data from the main memory (always a complete cache line):

$$\frac{n^2}{l} T_M$$

→ Loading data from the cache (cache hit because of “cache locality”):

$$(n^2 - \frac{n^2}{l}) T_C$$

- $T^{(1)}(n) = (2n^2 + n) T_A + (\frac{n^2}{l}) T_M + (n^2 - \frac{n^2}{l}) T_C$
- $T^{(1)}(n) = (2n^2 + n) T_A + \frac{n^2}{8} 180 T_A + 35(n^2 - \frac{n^2}{8}) T_A$
- $T^{(1)}(n) = (55,125 n^2 + n) T_A$
- *Cache – miss – ratio* = $(n^2/l)/n^2 = 1/l$

$T_M = 180 T_A$
$T_C = 35 T_A$
$l = 8$

1. Norm calculation of a matrix

c) norm1_col_wise()

- c) By switching the order of the two for loops in a), the algorithm now computes $\|A\|_1$. What time $T^{(2)}(n)$ does this algorithm take? Calculate the cache-miss-ratio.

Livedemo


1. Norm calculation of a matrix

c) norm1_col_wise()

```
double norm1_col_wise(double** const A, const int n) {
    double colsum = 0., max_norm = -1.;
    /* TODO Implement the one norm by switching
       the 2 for loops from the max norm */
    for (int j = 0; j < n; ++j) {
        colsum = 0.;
        for (int i = 0; i < n; ++i)
            colsum += abs(A[i][j]);
        if (colsum > max_norm)
            max_norm = colsum;
    }
    return max_norm;
}
```


1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()




1. Norm calculation of a matrix

c) norm1_col_wise()



1. Norm calculation of a matrix

c) norm1_col_wise()



1. Norm calculation of a matrix

c) Time for $T^{(2)}$

- Time for arithmetic operations does not change

- Case $n \cdot l < c$ (“Matrix fits into the cache”):

$$\rightarrow T^{(2b)}(n) = T^{(1)}(n) = (2n^2 + n) T_A + \frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C$$

$$\rightarrow \text{Cache} - \text{miss} - \text{ratio} = \frac{\frac{n^2}{l}}{n^2} = \frac{1}{l}$$

- Cases $n > c$ and $n \cdot l > c$ (“Matrix does not fit into the cache”):

- No cache hit: Every value has to be loaded from the main memory

- $T^{(2)}(n) = (2n^2 + n) T_A + n^2 T_M$

- $\text{Cache} - \text{miss} - \text{ratio} = \frac{n^2}{n^2} = 1$

1. Norm calculation of a matrix

d) norm1_row_wise()

- d) Design an algorithm to compute $\|A\|_1$ with consecutive memory access and calculate its runtime $T^{(3)}(n)$.

Hint: Use an auxiliary array for the column sums.

Livedemo

1. Norm calculation of a matrix

d) norm1_row_wise()

```
double norm1_row_wise(double** const A, const int n)
{
    /* TODO You need an auxiliary array
       for the column sums */
    double *colsums = new double[n];
    double max_norm = -1.;

    /* TODO The auxiliary array must be initialised */
    for (int i = 0; i < n; ++i)
        colsums[i] = 0;
```

Hint: $O(n)$ → We neglect this for T later.

...

1. Norm calculation of a matrix

d) norm1_row_wise()

...

```
/* TODO Now compute the column sums with  
consecutive memory access */
```

```
for (int i = 0; i < n; ++i)  
    for (int j = 0; j < n; ++j)  
        colsums[j] += abs(A[i][j]);
```

```
/* TODO Find the largest column sum */
```

```
for (int i = 0; i < n; ++i)  
    if (colsums[i] > max_norm)  
        max_norm = colsums[i];
```


Hint: $O(n) \rightarrow$ We neglect this for T later.

```
delete[] colsums;  
return max_norm;
```

```
}
```


1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()




1. Norm calculation of a matrix

d) norm1_row_wise()



1. Norm calculation of a matrix

d) norm1_row_wise()



1. Norm calculation of a matrix

d) Time for $T^{(3)}$

- Assumption: Auxiliary array is much bigger than the cache
- Time for arithmetic operations does not change
- Additional memory access to auxiliary array:

$$\frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C$$

Hint: We neglect the initialization of the auxiliary array and the comparison, because they are small compared to this term ($O(n)$). If you want to respect them you need to add $\frac{n}{l} T_M + \left(n - \frac{n}{l}\right) T_C$.

- $T^{(3)}(n) = T^{(1)} + \frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C =$

$$(2n^2 + n) T_A + \frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C + \frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C =$$

$$(2n^2 + n) T_A + 2 \left[\frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l}\right) T_C \right]$$

1. Norm calculation of a matrix

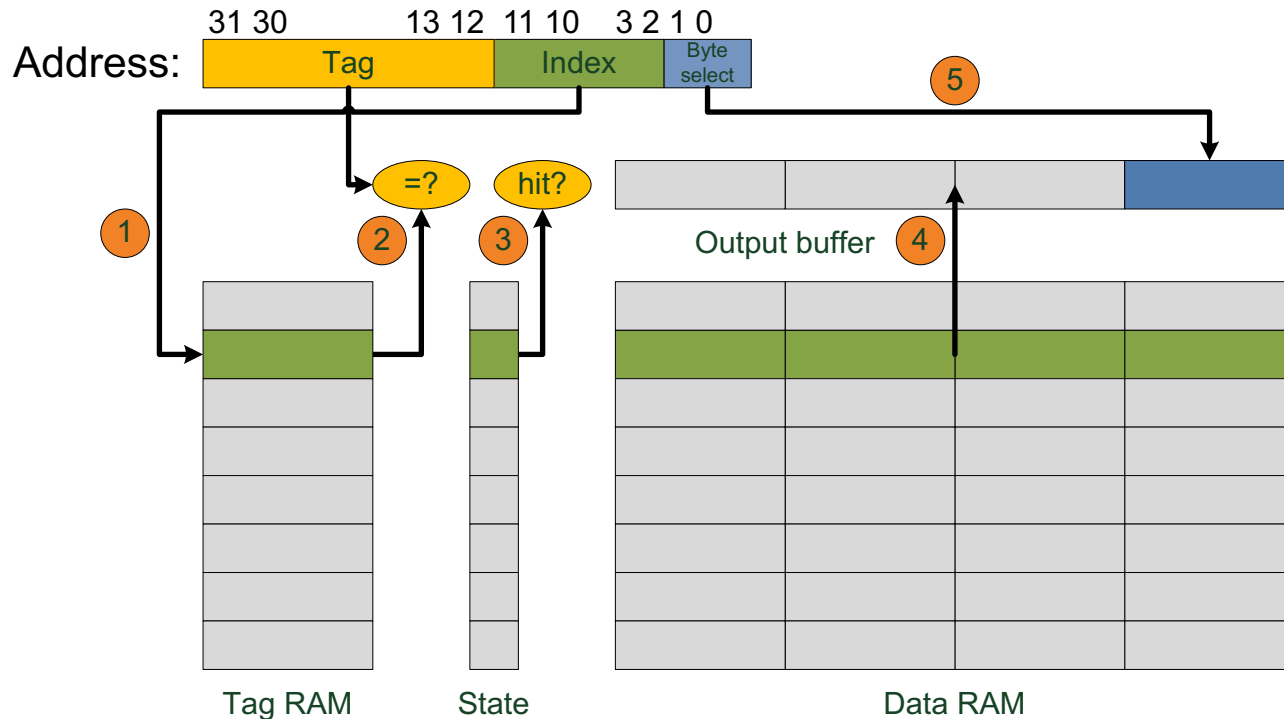
e) Cache mapping

- e) What is the worst case scenario regarding cache-miss-ratio for the algorithm of d) when we assume a direct-mapped cache?

Recap: Cache organization

■ Direct Mapped Cache:

→ Part of the address selects one entry in the tag RAM for comparison



1. Index selects one cache line
2. Check if selected tags are equal (no => miss)
3. Check if cache line is valid (no => miss)
4. Copy data to output buffer
5. Select required part from cacheline (if data is valid)

Exercise 02

Julian Miller | IT Center der RWTH Aachen University

1. Norm calculation of a matrix

e) Cache mapping

■ Worst case:


$\text{colsums}[i]$ and $A[x][i]$ have same index

$$(\text{colsums}[i] \equiv A[x][i] \bmod c)$$

■ $\text{colsums}[i] += A[x][i]$

→ LOAD of $A[x][i]$ evicts $\text{colsums}[i]$ from cache

→ LOAD/STORE of $\text{colsums}[i]$ has always cache miss



1. Norm calculation of a matrix

f) Calculate the Speed-up

f) Calculate the speed-up of the algorithm of d) with respect to the one of c).

$$\text{Hint: } \text{Speedup} = \frac{T^{(2)}(n)}{T^{(3)}(n)}$$

Assume that the following relations hold for an imaginary processor type:

$$T_M = 180 T_A$$

$$T_C = 35 T_A$$

$$l = 8.$$

(optional) Develop a block-version of the algorithm of d) and formulate its runtime $T^{(4)}(n)$ (cf. slide 302 f.).

1. Norm calculation of a matrix

f) Calculate the speed-up

- $T^{(2)}(n) = (2n^2 + n) T_A + n^2 T_M$

- **If the auxiliary array is much bigger than the cache**

- $T^{(3)}(n) = (2n^2 + n) T_A + 2 \left[\frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l} \right) T_C \right]$

- **In the blocked version the used part of the auxiliary array always fits into the cache**

- $T^{(4)}(n) = (2n^2 + n) T_A + \frac{n^2}{l} T_M + \left(n^2 - \frac{n^2}{l} \right) T_C + n^2 T_C$

1. Norm calculation of a matrix

f) Calculate the speed-up

- `norm1_col_wise()` : $T^{(2)}(n) = (182n^2 + n)T_A$
- `norm1_row_wise()` : $T^{(3)}(n) = (108.25n^2 + n)T_A$
- `norm1_blocking()` : $T^{(4)}(n) = (90.125n^2 + n)T_A$

■ Speed-up of the different Versions of `norm1_*()`


for big n




$$\rightarrow S(n) = \frac{T_{base}(n)}{T_{opt}(n)} = \lim_{n \rightarrow \infty} \left(\frac{a \cdot n^2 + n}{b \cdot n^2 + n} \right) = \lim_{n \rightarrow \infty} \left(\frac{a + \frac{1}{n}}{b + \frac{1}{n}} \right) = \frac{a}{b}$$

Algorithms	Speed-up
$T^{(2)} / T^{(3)}$	1.68
$T^{(2)} / T^{(4)}$	2.02
$T^{(3)} / T^{(4)}$	1.20


1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()



1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()

...


1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()




1. Norm calculation of a matrix (optional) norm1_blocking()



1. Norm calculation of a matrix (optional) norm1_blocking()



1. Norm calculation of a matrix (optional) norm1_blocking()



1. Norm calculation of a matrix

(optional) norm1_blocking()

```
double norm1_block(double** const A, const int n, const int l){
    double *colsums = new double[n]; double max_norm = -1.;

    /* TODO Copy the initialization of the auxiliary array and
    computation of the column sums. Implement the blocking with a
    surrounding for loop. */
    for (int c = 0; c < n; c += l) {
        int runto = std::min(c + l, n);
        for (int j = c; j < runto; ++j)
            colsums[j] = 0.;
        for (int i = 0; i < n; ++i)
            for (int j = c; j < runto; ++j)
                colsums[j] += abs(A[i][j]);
    }

    ...
}
```

1. Norm calculation of a matrix (optional) norm1_blocking()

...

```
/* TODO Find the largest column sum */  
    for (int i = 0; i < n; ++i)  
        if (colsums[i] > max_norm)  
            max_norm = colsums[i];  
  
    delete[] colsums;  
    return max_norm;  
}
```

Exercise Tasks

1. Norm calculation of a matrix
2. **Performance analysis tools**
3. Balance Metric

2. Performance analysis tools

Problem 2. Performance analysis tools

If you decided to solve Problem 1 by pen and paper, use the provided binaries in the binary directory in `exercise2.tar.gz`. Otherwise use your own code.

Problem 2.1. Profiling with gprof

You may run `make gprof` or execute `binary/norm.gprof.exe` and inspect the result with `gprof binary/norm.gprof.exe | less`

What is the time for the various algorithms according to gprof?

Try the target `gprof2`: `make gprof2` or execute `binary/norm.gprof2.exe` and inspect the result with `gprof binary/norm.gprof2.exe | less`

What is the difference of the outputs (You may have a look at the gprof targets in the Makefile)? What is the calculation time for the algorithms?

Livedemo

2.2 Hardware Counter

Problem 2.2. Hardware Performance Counter

Modern processor architectures have special-purpose registers to store the counts of hardware-related activities. These so called hardware performance counters can be used for low-level performance analysis or tuning. For accessing the hardware performance counters you can use the tool `likwid`¹ which is installed on the RWTH Compute Cluster.

Note 1: For this task you can use your own implementation of `a`), `c`), and `d`) from the exercise *norm calculation of a matrix* and build them with the Intel compiler (`make single` in the provided template) or use the provided binaries (`norm_max.exe`, `norm_1_col.exe`, and `norm_1_row.exe`).

Note 2: Use the special tuning node `login18-t.hpc.itc.rwth-aachen.de` for your measurements. Make sure that you are the only user on the system (e.g., with the `w` command) to obtain reliable results.

1. Examine the different programs with hardware counters using `likwid-perfctr`:

```
module load likwid
```

```
likwid-perfctr -C <processor id> -g <performance group> norm_max.exe
```

Use the performance group `MEM_DP` for your investigations.

2.2 Hardware Counter

- **HW performance counter: Special purpose registers for the count of HW events**
- **Can be used for low-level performance analysis or tuning**
- **The tool Likwid can be used**
 - Counters might be hard to interpret
 - Likwid uses derived metric (e.g., memory bandwidth, data volume, MFlop/s)
 - Use login18-t for your tests

Ignore the warning that you are not in the likwid group.
Using likwid like described should work nevertheless.
If not → mail to contact@hpc

2.2 Hardware Counter

1) Using likwid-perfctr

- Load the likwid module

```
$ module load likwid
```

- Pin the program to one processor and examine the hardware counters of this processor with likwid

```
$ likwid-perfctr -C <processor id> -g <performance group>  
norm_max.exe
```

- Use the performance group MEM_DP

- N=32768

```
→ $ likwid-perfctr -C 6 -g MEM_DP norm_max.exe 32768
```


2.2 Hardware Counter

1) Using likwid-perfctr

■ How does this look like?

→ Part I:
Counters

Group 1: MEM_DP

Event	Counter	Core 6
INSTR_RETIRED_ANY	FIXC0	4181301685
CPU_CLK_UNHALTED_CORE	FIXC1	13407652411
CPU_CLK_UNHALTED_REF	FIXC2	7798376880
PWR_PKG_ENERGY	PWR0	306.8376
PWR_DRAM_ENERGY	PWR3	83.8334
FP_ARITH_INST_RETIRED_128B_PACKED_DOUBLE	PMC0	327680
FP_ARITH_INST_RETIRED_SCALAR_DOUBLE	PMC1	1146938
FP_ARITH_INST_RETIRED_256B_PACKED_DOUBLE	PMC2	1342504960
FP_ARITH_INST_RETIRED_512B_PACKED_DOUBLE	PMC3	0
CAS_COUNT_RD	MBOX0C0	230072867
CAS_COUNT_WR	MBOX0C1	49141767
CAS_COUNT_RD	MBOX1C0	230072655
CAS_COUNT_WR	MBOX1C1	49142741
CAS_COUNT_RD	MBOX2C0	230070771
CAS_COUNT_WR	MBOX2C1	49142542
CAS_COUNT_RD	MBOX3C0	20616
CAS_COUNT_WR	MBOX3C1	21615
CAS_COUNT_RD	MBOX4C0	20943
CAS_COUNT_WR	MBOX4C1	21603
CAS_COUNT_RD	MBOX5C0	21311
CAS_COUNT_WR	MBOX5C1	21625

2.2 Hardware Counter

1) Using `likwid-perfctr`

■ How does this look like?

→ Part II: Metrics

Metric	Core 6
Runtime (RDTSC) [s]	5.3026
Runtime unhalted [s]	6.4029
Clock [MHz]	3600.1816
CPI	3.2066
Energy [J]	306.8376
Power [W]	57.8650
Energy DRAM [J]	83.8334
Power DRAM [W]	15.8097
DP MFLOP/s	1013.0461
AVX DP MFLOP/s	1012.7062
Packed MUOPS/s	253.2383
Scalar MUOPS/s	0.2163
Memory read bandwidth [MBytes/s]	8331.2911
Memory read data volume [GBytes]	44.1779
Memory write bandwidth [MBytes/s]	1780.1463
Memory write data volume [GBytes]	9.4395
Memory bandwidth [MBytes/s]	10111.4375
Memory data volume [GBytes]	53.6173
Operational intensity	0.1002

2.2 Hardware Counter

2. Analysis of the results. Likwid derives different metric from the counted events. Answer the following questions:

- Let be $n = 2^{15} = 32768$. Determine the read data volume D_{read} for the computation of $\|A\|_{\infty}$? What is reported by likwid? Why do these values differ?
- Which floating point performance (in MFLOP/s) is reported by likwid? Why is it not as high as reported by the norm application?
- Which memory bandwidth is reported by likwid? How can you assess whether the value is reasonable?

2.2 Hardware Counter

- Let be $n = 32768$. Determine the read data volume D_{read} for the computation of $\|A\|_{\infty}$

→ $D_{read} = n * n * 8 \text{ Byte} = 8.59 \text{ GB}$

- What is reported by likwid?

→ 44.18 GB

- Why do these values differ?

→ Code in main():

```
for (int i = 0; i < num_tests; ++i) {  
    n1 = norm_max(A, n),  
}
```

→ Computation is executed $\text{num_tests} = 5$ times.

→ $5 * 8.59 \text{ GB} = 42.95 \text{ GB}$

→ Note: Some compiler (e.g., clang) optimize the code such that the loop is only executed once (might lead to wrong interpretation of the performance)

2.2 Hardware Counter

- **Which floating point performance is reported by likwid?**

- 1012.87 Mflop/s

- **Why is it not as reported by the norm application?**

- Application reports 1650 Mflop/s

- Reason: likwid measures the complete application (including matrix allocation, initialization and error checking)

- Solution 1: Use likwid marker API to only measure the kernel

- Solution 2: Make more repetition (num_tests variable) to make kernel more dominant

- num_tests = 100: 1594.87 Mflops/s

2.2 Hardware Counter

- **Which memory bandwidth is reported by likwid?**

- 10110.85 MB/s

- **How can you assess whether the value is reasonable?**

- STREAM benchmark for one thread: 13842.9 MB/s (Triade)

- Not optimal, but reasonable

- Measured for complete application (including matrix allocation, initialization and error checking)

- Value for num_tests = 100: 12896.04 MB/s

Exercise Tasks

1. Norm calculation of a matrix
2. Performance analysis tools
3. **Balance Metric**

3. Performance Modeling

- **Observation:**
We reached a performance of 1.6 Gflop/s on a given architecture.
- **Question:**
Is that a good or a bad performance?
- **Answer:**
Can be given an adequate performance model.

3.1 Machine Balance

■ Reference machine: 2-socket Intel Skylake CPU

- 48 cores in total
- 2.1 GHz (base) clock frequency
- AVX-512 registers (512 bits)
- 2 operations per cycle (FMA)
- 2 FMA units per core
- Cache sizes: L1 → 32 KiB, L2 → 1 MiB, L3 → 33 MiB
- Sustainable main mem. bandwidth gained by Stream benchmark

See lecture
slide Part 4,
page 12 f.

■ Compute the machine balance B_m of this architecture (doubles).

■ Solution

$$\rightarrow B_m = \frac{b_s}{P_{max}}$$

b_s : achievable bandwidth over slowest data path [words/s]
 P_{max} : peak amount of floating-point operations per seconds [Flop/s]

3.1 Machine Balance

■ Achievable bandwidth [words/s]

$$\rightarrow b_s = \frac{STREAM\ BANDWIDTH}{8\ B} = \frac{170\ GB/s}{8\ B} = 21.25\ GWords/s$$

■ Peak amount of floating-point operations

$$\rightarrow P_{max} = \#Cores * Frequency * SIMD_OPs$$

$$\rightarrow SIMD_OPs = 2 * 2 * 8\ Flop = 32\ Flops$$

→ 2 operations (multiply + add)

→ 8 doubles per AVX-512 register (512 bit / 64 bit = 8, length of double: 8 Byte)

→ 2 FMAs

■ Solution

$$\rightarrow B_m = \frac{b_s}{P_{max}}$$

$$= \frac{\frac{170\ GB/s}{8\ B}}{48 \cdot 2.1\ GHz \cdot 32\ Flop} = \frac{21.25\ GWords/s}{3225.6\ GFlop/s} \approx 0.0067 \frac{Words}{Flop}$$

b_s : achievable bandwidth over slowest data path [words/s]

P_{max} : peak amount of floating-point operations per seconds [Flop/s]

3.1 Machine Balance

- With respect to the norm calculation of a matrix, which performance limits could be adapted?

- **Solution**

→ $B_m = \frac{b_s}{P_{max}}$

b_s : achievable bandwidth over slowest data path [words/s]

P_{max} : peak amount of floating-point operations per seconds [Flop/s]

→ The norm matrix computation is a serial application!

→ Previously computed P_{max} and b_s can never be reached

→ Use single-core values for new B_m :

See lecture
slide Part 4,
page 12 f.

$$B_m = \frac{\frac{14 \text{ GB/s}}{8 \text{ B}}}{1 \cdot 2 \cdot 1 \cdot 32 \text{ GFlop/s}} = \frac{1.75 \text{ GWords/s}}{67.2 \text{ GFlop/s}} \approx 0.026 \frac{\text{Words}}{\text{Flop}}$$

→ Use this performance limit in the following

3.2 Code Balance

See lecture
slide Part 4,
page 14 f.

- **Compute the code balance B_c of the $|A|_\infty$ computation.**
- **Assumptions**
 - Access to main memory is the slowest data path
 - Floating-point comparison ($a < b?$) = 1 is NOT a floating-point operation
 - Ignore the (statistical) sign flipping in the abs-function
- **Which elements are located in registers, caches and main memory?**

3.2 Code Balance

■ $|A|_{\infty}$ computation

1 FP operations

1 load
(A needs $32,768^2 \cdot 8$ B
= 8 GB in memory,
rowsum in register)

■ Solution

```
double norm_max(double** const A, const int n)
{
    double rowsum=0., max_norm=-1.;
    for (int i=0; i<n; ++i)
    {
        rowsum=0.;
        for (int j=0; j<n; ++j)
            rowsum += abs(A[i][j]);
        if (rowsum>max_norm)
            max_norm=rowsum;
    }
    return max_norm;
}
```

Ignore
comparison

```
double abs(double x) {
    return x<0 ? -x : x;
}
```

→ Code balance: $B_c = \frac{\text{data transfers}(\text{LOAD,STORE})}{\text{arithmetic operations}} \frac{[\text{Words}]}{[\text{Flop}]}$

$$\rightarrow B_c = \frac{1 \text{ Word}}{1 \text{ Flop}}$$

3.3 Lightspeed

■ Compute the (relative) lightspeed l of the $|A|_\infty$ computation on an Intel Skylake architecture

- How do you interpret this value?
- In general, how can a lightspeed value be improved by the application developer?

See lecture
slide Part 4,
page 17 f.

■ Solution

- Lightspeed $l = \min(1, \frac{B_m}{B_c})$ with $B_m = 0.026 \frac{\text{Words}}{\text{Flop}}$, $B_c = 1 \frac{\text{Word}}{\text{Flop}}$
- $l = \min(1, \frac{0.026}{1}) = 0.026$
- 2.6 % of machine's peak performance possible for this code
- Machine balance cannot be influenced; only code balance can be influenced by improving the transfer/computation ratio (e.g. by better cache usage)

3.3 Lightspeed

- Compute the lightspeed P for absolute performance in GFlop/s!

- Solution

$$\rightarrow P = l \cdot P_{max} = \min\left(P_{max}, \frac{b_s}{B_c}\right) \quad \text{with} \quad b_s = 1.75 \frac{G \text{ Words}}{s}, \quad B_c = 1 \frac{\text{Word}}{\text{Flop}}$$

$$\begin{aligned} \rightarrow P &= \min(67.2 \text{ GFlop/s}, \frac{1.75 \text{ G words/s}}{1 \text{ word/Flop}}) \\ &= \min(67.2 \text{ GFlop/s}, 1.75 \text{ GFlop/s}) = 1.75 \text{ GFlop/s} \end{aligned}$$

3.3 Lightspeed

- **Compare the measured performance [GFlop/s] (from the previously-done likwid VIEW run) with the computed theoretical performance**
 - How close is the code to what it can reach at maximum?
 - Why can there be a difference between theoretical and experimentally-gathered results?
- **Solution**
 - Theoretical performance limit: $P = 1.75 \text{ GFlop/s}$
 - Performance analysis: $\sim 1.6 \text{ GFlops/s}$
 - Balance metric uses certain assumptions (that usually cannot hold completely), e.g.
 - Data transfer & arithmetical operations overlap perfectly
 - Perfect use of FMA (here no multiplications needed, not relevant if memory bound)
 - Only the slowest data path is/ can be modeled (others are infinitely fast)
 - Latency is ignored

Real code achieves 91% of theoretical peak performance of the code