

# Exercise 3

Introduction to High-Performance Computing  
WS 2019/20

Simon Schwitanski

Christian Terboven

[contact@hpc.rwth-aachen.de](mailto:contact@hpc.rwth-aachen.de)

1. **Scaling behavior**
2. **Cache coherence**
3. **Key figures of networks**
4. **Task dependency graphs**

# 1. Scaling behavior

## 1.1. Strong scaling



- **Given: parallel execution time with 4 processes = 12h**  
**10% is serial processing**

→  $T(4) = 12h$ ,  $s_{4procs} = 0.1$  ( $s$  for 4 processors!)

→ Serial processing time is  $0.1 \cdot 12h = 1.2h$

→ Parallel processing time is  $10.8h$



- a) **Calculate compute time on 16 parallel processes!**

→ 
$$T(16) = s_{4procs} \cdot T(4) + p_{4procs} \cdot \frac{T(4)}{4} = 1.2h + \frac{10.8h}{4} = 3.9h$$

(Caution: Baseline was given with 4 procs., so we divide by 4 instead of 16!)

- b) **How long takes serial execution? What is the sequential portion?**

→ Serial execution:  $T(1) = s_{4procs} \cdot T(4) + 4 \cdot p_{4procs} \cdot T(4) = 44.4h$

→ Sequential portion: 
$$s_{1proc} = \frac{s_{4procs} \cdot T(4)}{T(1)} = \frac{1.2h}{44.4h} = 2.7 \% = 0.027$$

# 1. Scaling behavior

## 1.1. Strong scaling



- c) How many processes are needed to get a 2-hour execution time?  
What is the lower bound of parallel calculation time?**

→ Intuitively:

→ 2h total runtime - 1.2h sequential runtime = 0.8h parallel runtime

→ 43.2h parallel workload / (0.8h runtime per process) = 54 processes

→ Formally:

→  $T(N) = 2h$

→  $T(N) = s \cdot T(1) + \frac{p \cdot T(1)}{N} \Rightarrow N = \frac{p \cdot T(1)}{T(N) - s \cdot T(1)} = \frac{43.2h}{0.8h} = 54$

→ Lower bound of parallel calculation time  $T(N)$

→ For  $N \rightarrow \infty$ : runtime  $\rightarrow$  sequential runtime = 1.2h

→  $\lim_{N \rightarrow \infty} T(N) = s \cdot T(1) + \frac{p \cdot T(1)}{N} = s \cdot T(1) = 1.2h$

# 1. Scaling behavior


## 1.1. Strong scaling



### d) Speedup based on serial runtime

$$\lim_{N \rightarrow \infty} S_p(N) = \lim_{N \rightarrow \infty} \frac{1}{s + \frac{1-s}{N}} = \frac{1}{s}$$

$$\rightarrow 1/0.027 = 37$$



# 1. Scaling behavior


## 1.1. Strong scaling



### d) Speedup based on parallel execution with 4 processes

$$\lim_{N \rightarrow \infty} S_p(N) = \lim_{N \rightarrow \infty} \frac{1}{s + \frac{1-s}{N}} = \frac{1}{s}$$

$$\rightarrow 1/0.1 = 10$$



# 1. Scaling behavior

## 1.2. Weak scaling



- **What parameter of a weather simulation could you adjust to have constant parallel computation time when you increase parallel computation capability?**
  - Time resolution
  - Spatial resolution
  - Size of the region

# 1. Scaling behavior


## 1.2. Weak scaling



- The basic strategy with weak scaling is to replicate workload for each process → fixed problem size per process


- Higher time resolution results in more iteration cycles:

→ If you increase the number of iteration cycles, you also add more calls to the sequential *update\_grid* (red part) function (example: doubling iteration cycles)



→ If you then increase the number of processors, then you decrease the time of each iteration cycle

→ **But:** The sequential portion of the calculation gets bigger which contradicts Gustafson's assumption of having a constant sequential portion






# 1. Scaling behavior


## 1.2. Weak scaling

- **Higher spatial resolution and increasing the size of the region both result in more grid points.**

- Parallel workload increases with number of processes (e.g. doubling number of processes requires doubling number of grid points for same problem size per proc.)
- Time per iteration cycle stays constant with increasing number of processes and no further calls to *update\_grid* are added (number of timesteps is fixed)



1. Scaling behavior
2. **Cache coherence**
3. Key figures of networks
4. Task dependency graphs



- After time step 3, cores see different values for X
- Depending on which cache writes back X, value might be stale


- **When switching to multicore processors, the concept of caching needs some considerations:**
  - Caches store
    - Private data only used by a single core
    - Data that is shared by multiple cores
  - Sharing data inside a common cache
    - Reduces memory latency on more than one core
    - Increases the effective bandwidth of shared data
  - Cache coherence problem
    - Modifications of cached data must be visible to all cores

## ■ Directory based

→ Sharing status of a block of physical memory is kept in only one place, the directory


## ■ Snooping

→ Caches are connected via a broadcast medium and snoop(monitor) on that medium for memory blocks that they currently store



■ **Cache controller snoops on the medium and take relevant action to ensure cache coherence:**

- Invalidate
- Update
- Supply value



■ **All recent MP systems use write-invalidate**

## ■ Write-through

- Data written to cache is directly written to memory
- Get most recent value from main memory

## ■ Write-back

- Cache lines are not directly written to main memory but are flagged as dirty and written back later on
- All cores check their caches for addresses placed on the memory bus
- If core has most recent copy of requested cache block it provides it in response to a read request

## ■ Write-back vs. Write-through

- Write-back needs lower memory bandwidth
- Most recent multiprocessors use write-back

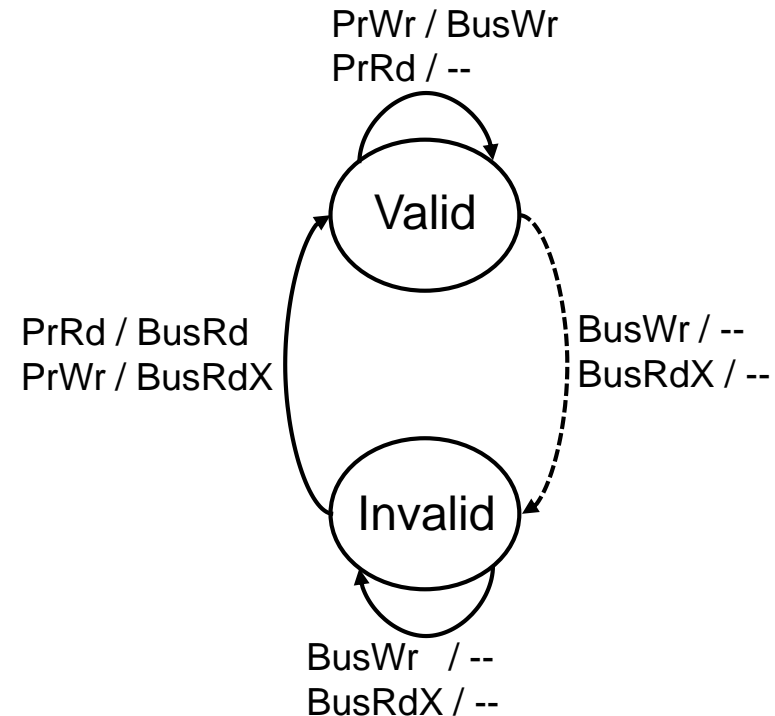


- Write-through snooping protocol based on block invalidation
- Each block of memory has one of the following states

→ Valid: block is present in cache,  
all copies of block are identical  
to copy in main memory

→ Invalid: block contains invalid  
data, needs the most updated  
copy from owner

→ PrWr on invalid data  
requires a first a read



PrRd: Processor read

PrWr: Processor write

BusRd (bus read): read request for a block

BusRdX (bus read exclusive): read block and invalidate other copies

BusWr: write a block to memory and invalidate other copies

# Cache coherence protocols

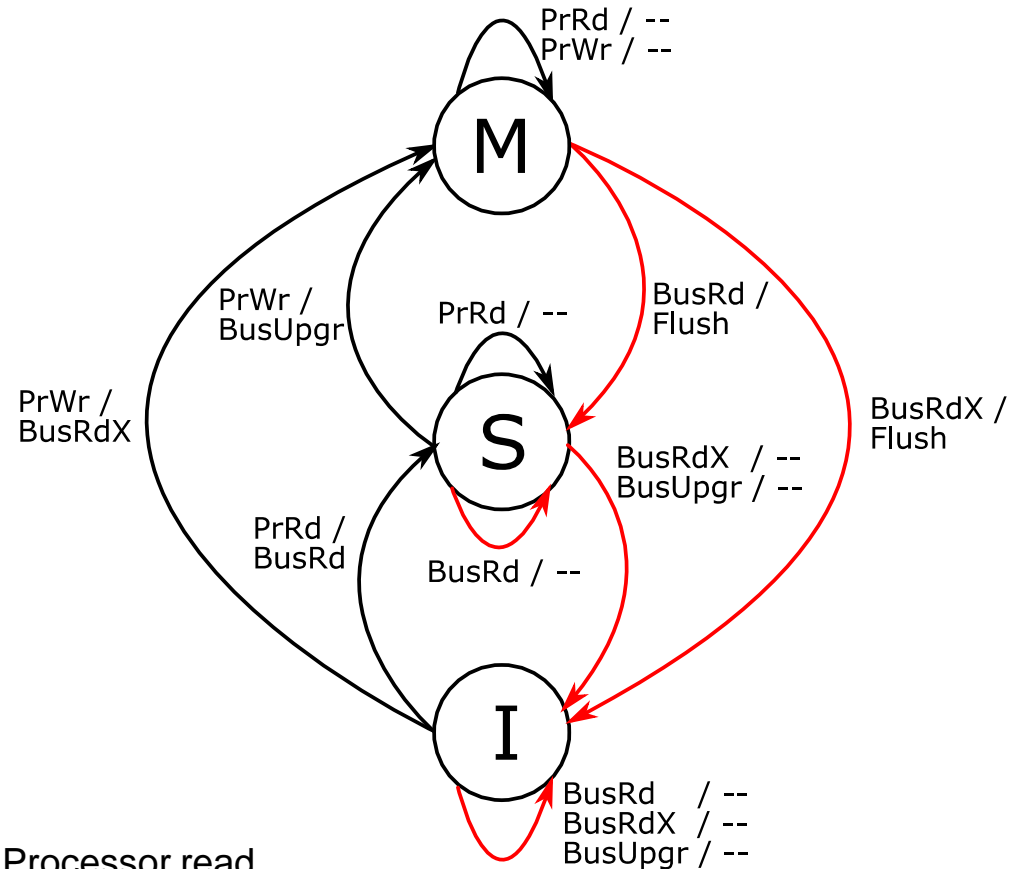
## Example: MSI invalidate protocol



- Write-back snooping protocol based on block invalidation

- Each block of memory has one of the following states

- Modified: exactly one cache (owner) holds a valid copy, memory is stale
- Shared: zero or more other caches hold valid copy
- Invalid: block contains invalid data, needs the most updated copy from owner



PrRd: Processor read

PrWr: Processor write

BusRd (bus read): read request for a block

BusRdX (bus read exclusive): read block and invalidate other copies

BusUpgr: invalidate other copies

Flush: Write cache block back to main memory

## 2. Cache Coherence

### 2.1 Applying the MSI Protocol



t	Local Request	P1	P2	P3	Gen. Bus Request	Data Supplier
0	Initially	-	-	-	-	-
1	R1	S	-	-	BusRd	Memory
2	W1	M	-	-	BusUpgr	-
3	R3	S	-	S	BusRd	P1's cache (flush)
4	W3	I	-	M	BusUpgr	-
5	R3	I	-	M	-	-
6	R2	I	S	S	BusRd	P3's cache (flush)

## 2. Cache Coherence

### 2.1 Applying the MSI Protocol



t	Local Request	P1	P2	P3	Gen. Bus Request	Data Supplier
0	Initially	-	-	-	-	-
1	W2	-	M	-	BusRdX	Memory
2	R2	-	M	-	-	-
3	R1	S	S	-	BusRd	P2's cache (flush)
4	W2	I	M	-	BusUpgr	-
5	R1	S	S	-	BusRd	P2's cache (flush)
6	R3	S	S	S	BusRd	P1/P2's cache (flush)
7	W1	M	I	I	BusUpgr	-
8	R1	M	I	I	-	-
9	W2	I	M	I	BusRdX	P1's cache
10	R3	I	S	S	BusRd	P2's cache

## 2. Cache Coherence

### 2.2 Adaptation of the MSI Protocol



#### ■ Add additional state “Exclusive”

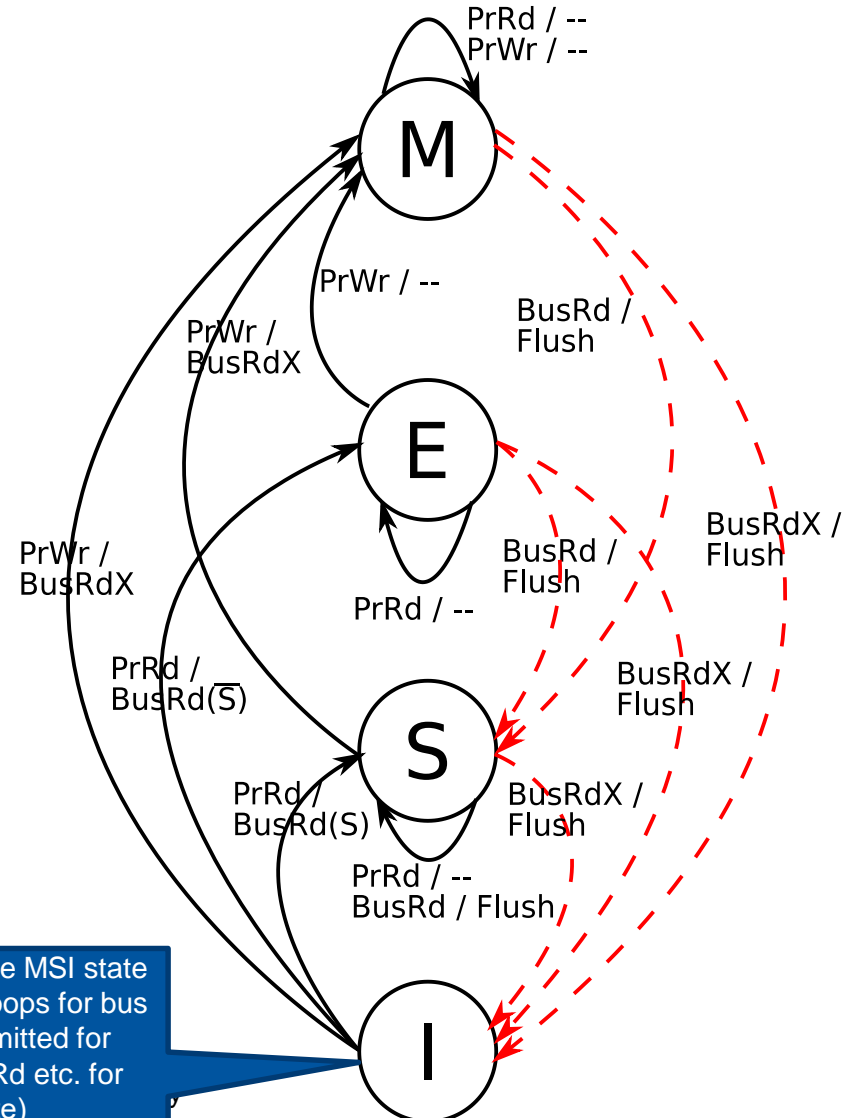
- Cache line has recently been read from memory but not yet modified. It does not reside in any other cache.

#### ■ Additionally: Shared signal “S”, determines on BusRd request if any other processor holds the cache block

- BusRd( $S$ ): As BusRd, in addition: There is a copy of the corresponding cache block on at least one other processor.
- BusRd( $\bar{S}$ ): As BusRd, in addition: There is no copy of the corresponding cache block on any other processor.
- If we write just “BusRd”, we do not care about the shared signal “S”.

Note: Compared to the MSI state machine, some self loops for bus transactions are omitted for simplicity (e.g. BusRd etc. for “Invalid” state)

#### MESI Protocol State Machine



## 2. Cache Coherence

### 2.2 Adaptation of the MSI Protocol



#### ■ MSI Protocol

t	Local Request	P1	P2	P3	Gen. Bus Request	Data Supplier
0	Initially	-	-	-	-	-
1	R1	S	-	-	BusRd	Memory
2	W1	M	-	-	BusUpgr	-
3	R3	S	-	S	BusRd	P1's cache (flush)

#### ■ MESI Protocol

t	Local Request	P1	P2	P3	Gen. Bus Request	Data Supplier
0	Initially	-	-	-	-	-
1	R1	E	-	-	BusRd	Memory
2	W1	M	-	-	-	-
3	R3	S	-	S	BusRd	P1's cache (flush)

1. Scaling behavior
2. Cache coherence
3. **Key figures of networks**
4. Task dependency graphs

## 3. Key figures of networks

### 3.1 Fully connected network



#### ■ Given: fully connected network (with N nodes)

##### a) Number of edges?

→ Standard handshaking problem: each node is connected with all others.

Count each edge only once: *number of edges* =  $\frac{N(N-1)}{2}$

##### b) Edge connectivity?

→ You need to remove all edges to one node to split the network, thus edge connectivity is N-1

##### c) Diameter? → 1



## 3. Key figures of networks

### 3.1 Fully connected network



#### d) Bisection bandwidth?

→ After bisecting the fully connected network with  $N$  nodes you have two fully connected networks with  $N/2$  nodes in case of even  $N$

→ **Case 1:**  $N$  even: Two equal halves of size  $\frac{N}{2}$

$$B_b = \frac{N(N-1)}{2} - 2 \frac{\frac{N}{2} \left( \frac{N}{2} - 1 \right)}{2} = \frac{N^2}{2} - \frac{N}{2} - \frac{N^2}{4} + \frac{N}{2} = \frac{N^2}{4}$$

→ If  $N$  is odd: One network with  $\frac{N+1}{2}$  nodes and one network with  $\frac{N-1}{2}$  nodes

→ **Case 2:**  $N$  odd: Network with  $\frac{N+1}{2}$  nodes and network with  $\frac{N-1}{2}$  nodes

$$\begin{aligned} B_b &= \frac{N(N-1)}{2} - \frac{\frac{N+1}{2} \left( \frac{N+1}{2} - 1 \right)}{2} - \frac{\frac{N-1}{2} \left( \frac{N-1}{2} - 1 \right)}{2} \\ &= \frac{1}{2} \left( N(N-1) - \frac{(N+1)(N-1)}{4} - \frac{N-1}{2} \left( \frac{N-1}{2} - 1 \right) \right) \\ &= \frac{1}{2} \left( N^2 - N - \frac{N^2 - 1}{4} - \frac{N^2 - 2N + 1}{4} + \frac{N-1}{2} \right) \\ &= \frac{1}{2} \left( N^2 - \frac{N^2}{4} - \frac{N^2}{4} - N + \frac{N}{2} + \frac{N}{2} + \frac{1}{4} - \frac{1}{4} - \frac{1}{2} \right) = \frac{1}{2} \left( \frac{N^2}{2} - \frac{1}{2} \right) = \frac{N^2}{4} - \frac{1}{4} \end{aligned}$$

## 3. Key figures of networks

### 3.1 Fully connected network



- **Combining both cases: The bisection bandwidth for a fully connected network with  $N$  nodes and bandwidth 1 on each link is**

$$B_b = \lfloor \frac{N^2}{4} \rfloor.$$

## 3. Key figures of networks

### 3.2 LogP Model



#### ■ Sketch a LogP model diagram for the global reduction-to-all communication operation

→ Given:  $P$  processes = 8, vectors of size  $n$ , binary reduce operation, computation time  $T_c(n) = 1$ , network transmission time  $T_n(n) = 1$

a) Cyclic reduction

b) Reduce-and-broadcast

→ Binary tree

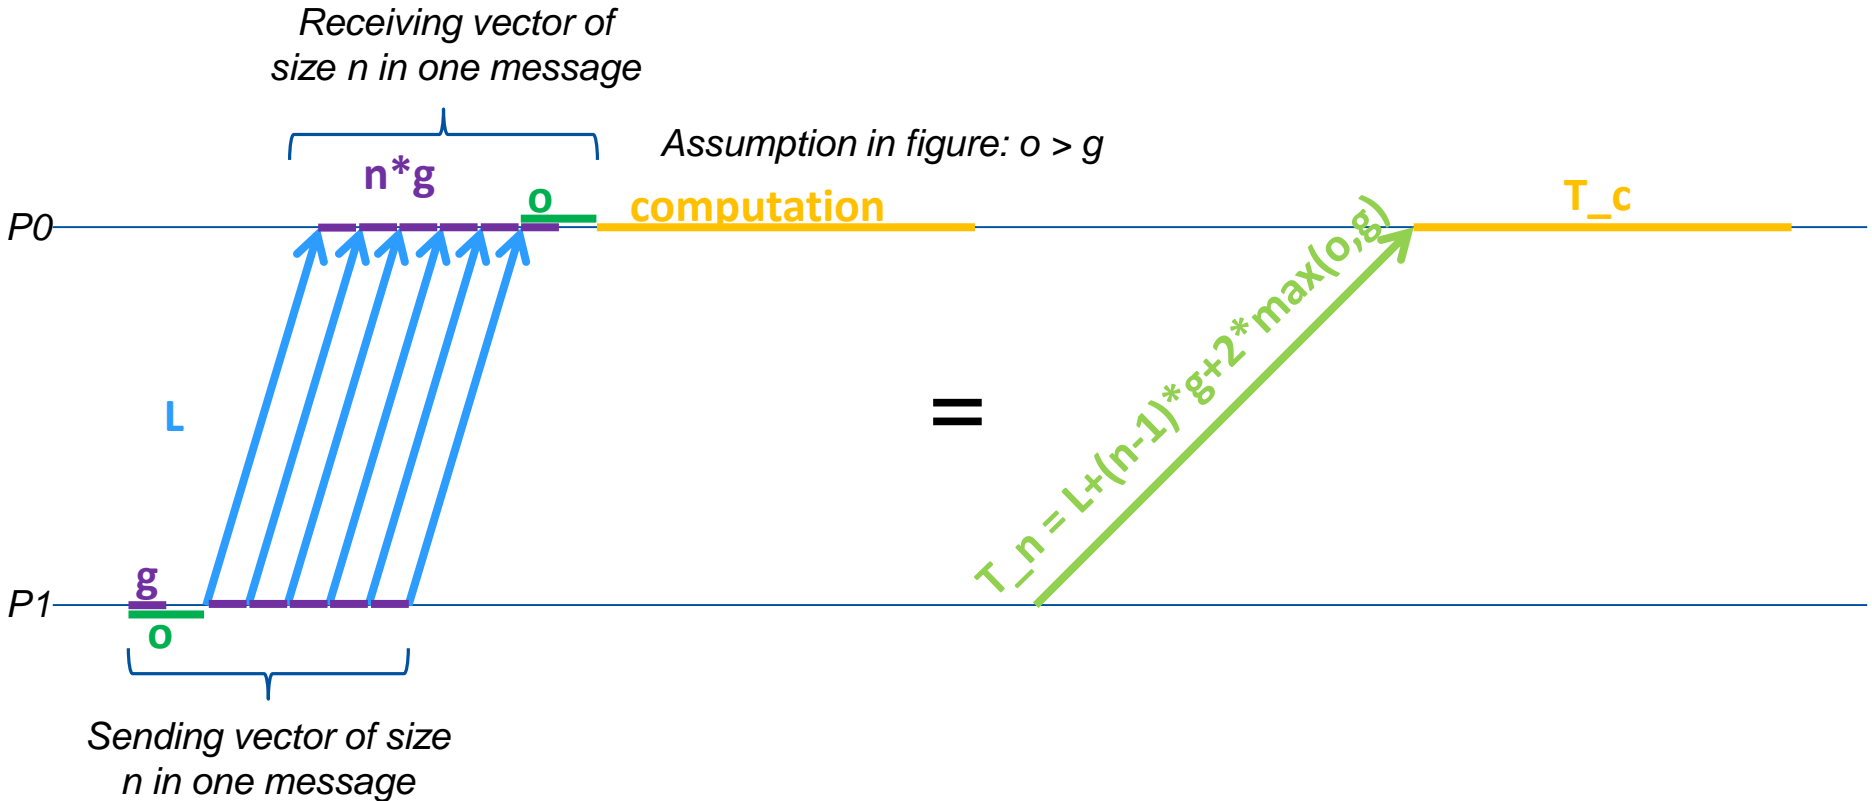
→ Binomial tree

#### ■ Compare the total runtime

#### ■ What potential effect has streaming (communicate calculated values) for large vectors?

# 3. Key figures of networks

## 3.2 LogP Model



Assumptions on the following slides:



$$T_c = 1, T_n = 1$$

# 3. Key figures of networks

## 3.2 LogP Model



■ **Cyclic reduction:**  $(P - 1) (T_c + T_n) = 7 \cdot 2 = 14$



# 3. Key figures of networks

## 3.2 LogP Model





### Binary tree:

→ reduction =  $4 (T_c + T_n) = 4 \cdot 2 = 8$

→ broadcast =  $4 T_n = 4 \cdot 1 = 4$

→ reduce – to – all =  $8 + 4 = 12$



*Assumption: communication & computation need roughly the same time*

→ Symmetric communication scheme

# 3. Key figures of networks


## 3.2 LogP Model



### ■ Binomial tree construction

→  $B_0$ : Binomial tree of order 0 is a single node

→  $B_k$ : Binomial tree of order  $k$  has a root node with  $k$  children, namely the binomial trees of orders  $0, 1, 2, \dots, k-1, k-2$



# 3. Key figures of networks

## 3.2 LogP Model





### ■ Binomial tree:

→  $\text{reduction} = \lceil \lg(P) \rceil (T_c + T_n) = 3 \cdot 2 = 6$

→  $\text{broadcast} = \lceil \lg(P) \rceil (T_n) = 3 \cdot 1 = 3$

→  $\text{reduce - to - all} = 6 + 3 = 9$



*Assumption: communication & computation need roughly the same time*

→ Better on broadcast and on reduction

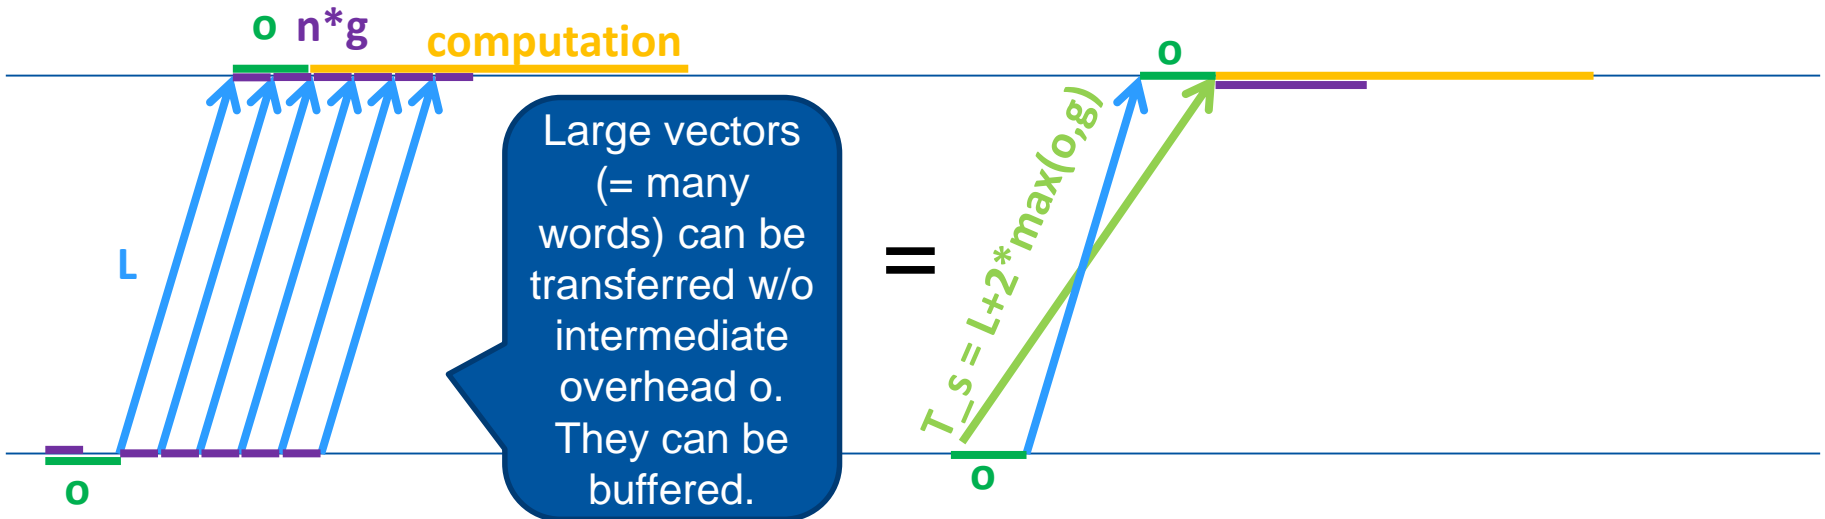


# 3. Key figures of networks

## 3.2 LogP Model



- c) Effect of streaming for large vectors
- For large transfers, LogP model is not the perfect model. Transmission is done by DMA controller, CPU is busy (o) just to start transfer.




- What effect does this have on the reduce-and-broadcast pattern?

→ In the reduce steps,  $T_n$  can be replaced by  $T_s$  (due to overlapping computation and communication)


1. Scaling behavior
2. Cache coherence
3. Key figures of networks
4. **Task dependency graphs**

### 3. Task dependency graphs


- Determine: total work, critical path length, average concurrency



Total work: 55  
Critical path length: 29  
Avg. concurrency: ~1.9



Total work: 69  
Critical path length: 39  
Avg. concurrency: ~1.8



Total work: 95  
Critical path length: 41  
Avg. concurrency: ~2.3