

# Machine Learning - Exercise 4

## Companion Slides

Ali Athar Sabarinath Mahadevan

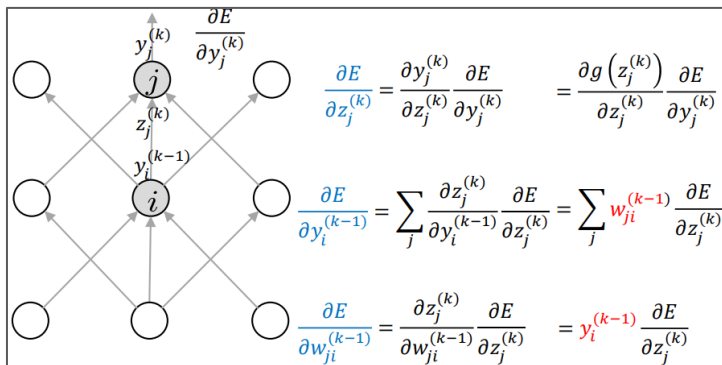
December 6, 2018



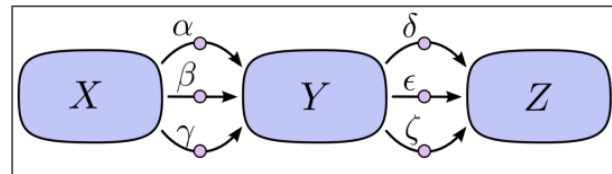
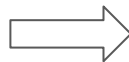
Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Exercise Goal



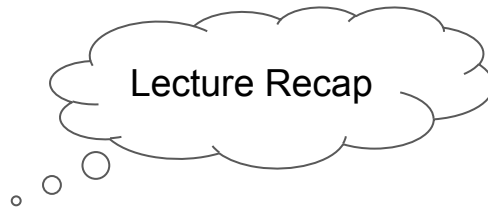
Backpropagation for fixed network



General backpropagation with computational graphs

This exercise is about

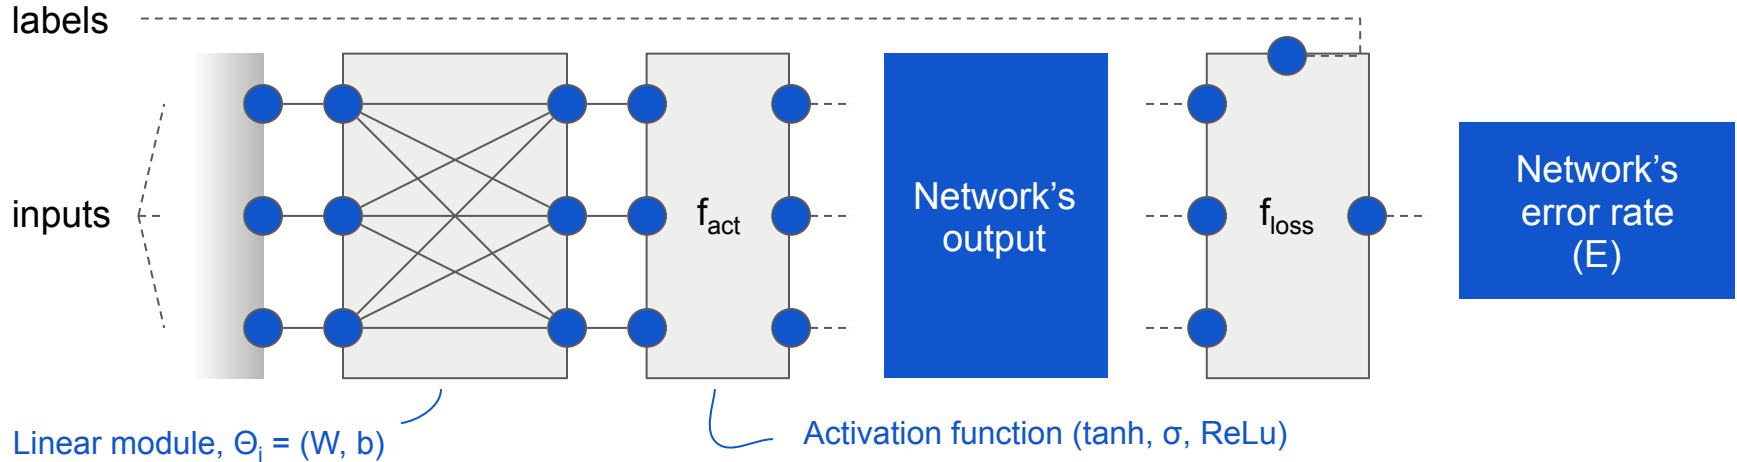
- ▶ Understanding backpropagation, deriving formulas, optimizing them
- ▶ Implement simple neural network framework yourself
- ▶ Digit recognition



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Recap: Neural Networks



- ▶ Training data (inputs)  $X = \{x_i\}_{i=1..N}$  with  $x_i \in \mathbb{I}$ ,  $N$  the batch size
- ▶ Training labels  $T = \{t_i\}_{i=1..N}$  with  $x_i \in \mathbb{O}$
- ▶ Network is a parametrized, (sub-)differentiable function  $F(X, \Theta) : \mathbb{I} \times \mathbb{P} \rightarrow \mathbb{O}$ 
  - ▶ e.g.,  $\mathbb{O} = \mathbb{R}^{\text{Dim}}$  (regression),  $\mathbb{O} = [0, 1]^{\text{Dim}}$  (prob. classification)
- ▶ Loss (criterion)  $\mathbf{L}(T, F(X, \Theta)) : \mathbb{O} \times \mathbb{O} \rightarrow \mathbb{R}$ , put on top of output to measure performance
  - ▶ find optimal parameters:  $\Theta^* = \text{argmin}_{\Theta} \mathbf{L}(T, F(X, \Theta))$

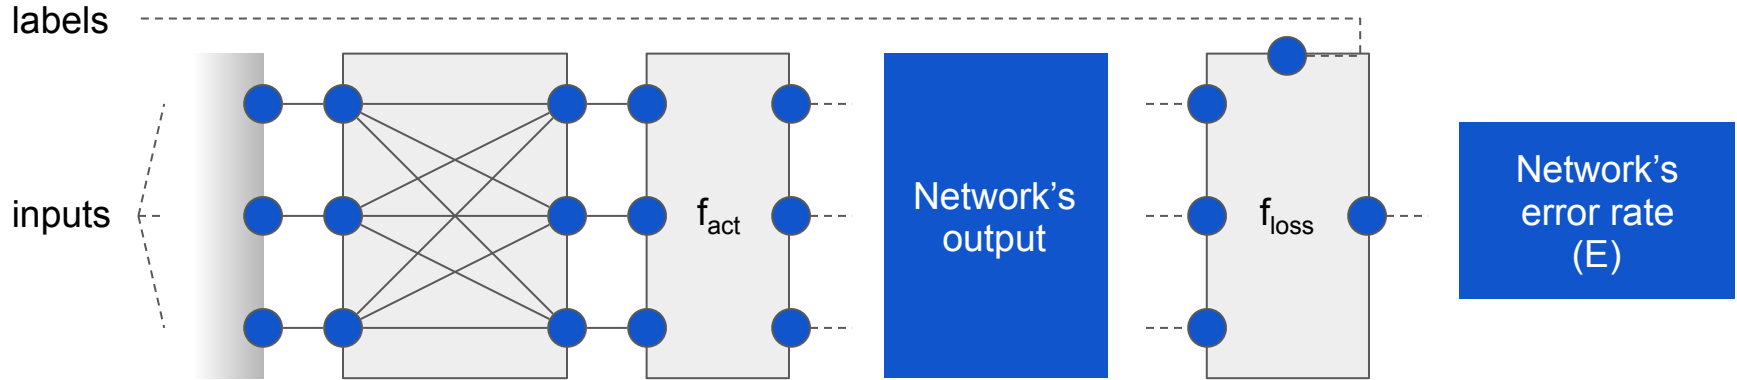
Parameters



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Recap: Backpropagation

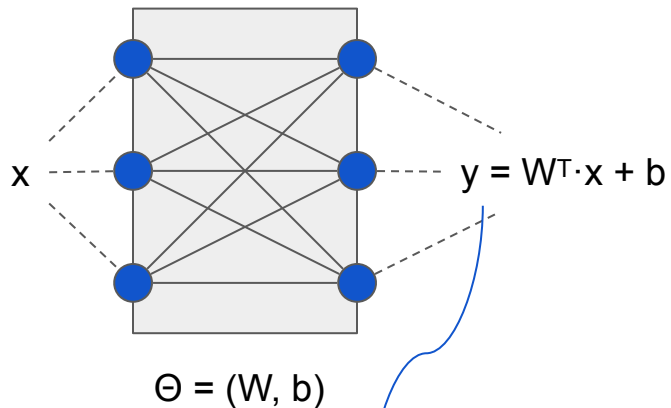


- ▶ Optimize towards lower error rate, i.e., lower  $E$ 
  - ▶ Take derivative of  $E$  with respect to each module's parameters, follow gradient
  - ▶ Example: Gradient Descent:  $\Theta = \Theta - \lambda * D_{\Theta}(E(x))$
  - ▶  $D_{\Theta}(E(x)) = D_{\Theta}(E)$  for brevity
- ▶ How to calculate  $D_{\Theta}(E)$ 
  - ▶ Reverse order of modules
  - ▶ Module gets  $D_{out}(E)$ , calculates  $D_{\Theta}(E)$ , passes  $D_{in}(E)$  to next module

Derivative w.r.t. module's parameters  $\Theta$  at point  $x$   
Learning rate



## Example: Linear/Fully Connected Module



Element-wise:

$$y = \left[ \sum_{i=1}^{N_{in}} (w_{ij} \cdot x_i) + b_j \right]_j$$



Without batching

Given: Derivative with respect to output  $\frac{dE(x)}{dy}$

Calculate:

- Derivatives with respect to parameters  $\Theta$

$$\frac{dE(x)}{dw_{ij}} = \frac{dE(x)}{dy} \cdot \frac{dy}{dw_{ij}} = \frac{dE(x)}{dy} \cdot x_i = \left[ x \cdot \frac{dE}{dy} \right]_{ij}$$

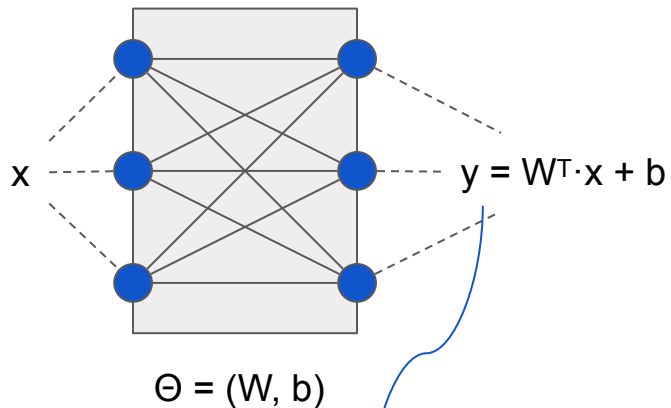
$$\frac{dE(x)}{db} = \frac{dE(x)}{dy} \cdot \frac{dy}{db} = \frac{dE(x)}{dy} \cdot I = \frac{dE(x)}{dy}$$



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

## Example: Linear/Fully Connected Module



Element-wise:

$$y = \left[ \sum_{i=1}^{N_{in}} (w_{ij} \cdot x_i) + b_j \right]_j$$



Without batching

Given: Derivative with respect to output  $\frac{dE(x)}{dy}$

Calculate:

- ▶ Derivative with respect to input

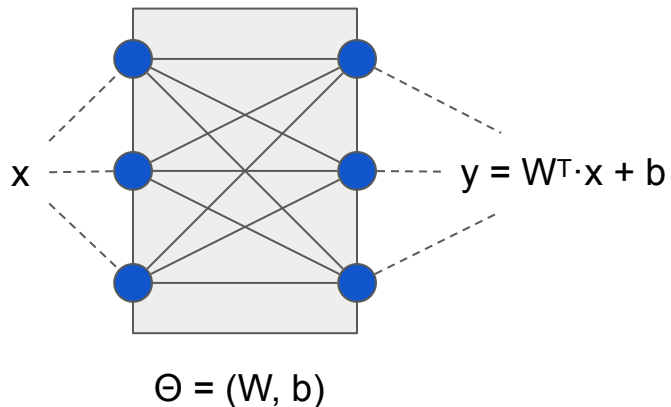
$$\frac{dE(x)}{dx} = \frac{dE(x)}{dy} \cdot \frac{dy}{dx} = \frac{dE(x)}{dy} \cdot W^T$$



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Example: Linear/Fully Connected Module



Putting it together:

```
fprop(x):  
    cache.x = x  
    return WT*x + b
```

run training data through  
(forwards)

```
bprop(dE):  
    dW = cache.x * dE  
    db = dE  
    return dE * W
```

run gradients through  
(backwards)

```
update(rate):  
    W = W - rate*dW  
    b = b - rate*dbT
```

update the parameters  
(grad. descent)



Without batching



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

# Mini-Batching

- ▶ Batch learning
  - ▶ All training samples processed at once, parameters updated once at the end
  - ▶ Stable, well understood, many acceleration techniques, but slow
- ▶ Stochastic learning
  - ▶ Each training sample separately, parameters updated at each step
  - ▶ Noisy (though may lead to better results), fast
- ▶ Mini-batching
  - ▶ Middle ground, batches of data processed, bundled updates
  - ▶ Combine advantages, reduce drawbacks
- ▶ Example
  - ▶ Linear Module  $f$  with input dimension  $N_{in}$  and output dimension  $N_{out}$ , batch size  $n$

$$f(x) = W^T \cdot \begin{bmatrix} | & | & & | \\ x^1 & x^2 & \dots & x^n \\ | & | & & | \end{bmatrix} + \begin{bmatrix} | & | & & | \\ b & b & \dots & b \\ | & | & & | \end{bmatrix}$$

broadcast (i.e., repeat)  $b$

mini-batch matrix



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY



# Batching Update Rule

- ▶ (Mini-)Batch learning

- ▶ Multiple samples processed at once
- ▶ Calculate gradient for each sample, but don't update the parameters
- ▶ After processing the batch, update using a sum of all gradients
- ▶ Learning rate has to be adapted, e.g., divide E by batch size

- ▶ Example: Gradient Descent  $\Theta = \Theta - \lambda \cdot \sum_{k=1}^N D_{\Theta}(E(x^k))$

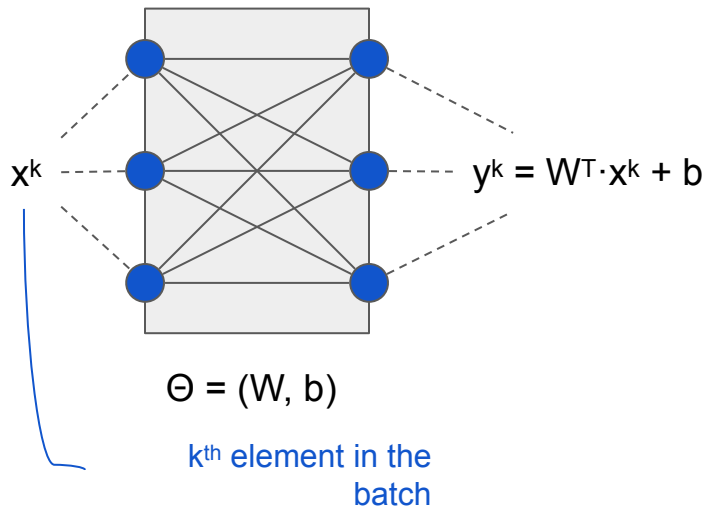
Derivative of E w.r.t parameters  $\Theta$  at point  $x^k$

- ▶ To make things easier, we write

$$\sum_{k=1}^N D_{\Theta}(E(x^k)) = D_{\Theta}(E(x)) = D_{\Theta}(E)$$



## Example: Linear/Fully Connected Module - Batching



Deriv. w.r.t outputs  
assumed to be  
given row-wise:

$$\frac{dE(x)}{dy} = \begin{bmatrix} - & \frac{dE(x^1)}{dy^1} & - \\ & \vdots & \\ - & \frac{dE(x^n)}{dy^n} & - \end{bmatrix}$$

Given: Derivatives with respect to outputs  $\frac{dE(x^k)}{dy^k}$

Calculate:

Plural!

- Derivatives with respect to parameters  $\Theta$

$$\begin{aligned} \frac{dE(x)}{dw_{ij}} &= \sum_{k=1}^N \left( \frac{dE(x^k)}{dy^k} \cdot \frac{dy^k}{dw_{ij}} \right) = \left[ \sum_{k=1}^N \left( x^k \cdot \frac{dE(x^k)}{dy^k} \right) \right]_{ij} \\ &= \left[ x \cdot \frac{dE(x)}{dy} \right]_{ij} \end{aligned}$$

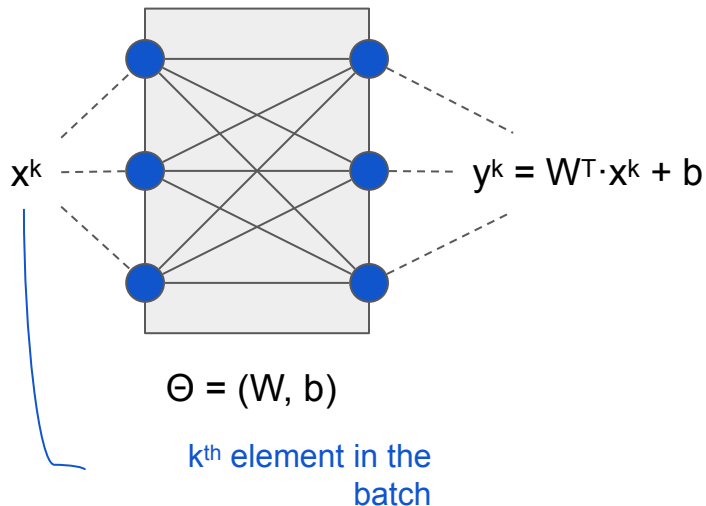
$$\begin{aligned} \frac{dE(x)}{db} &= \sum_{k=1}^N \left( \frac{dE(x^k)}{dy^k} \cdot \frac{dy^k}{db} \right) = \sum_{k=1}^N \left( \frac{dE(x^k)}{dy^k} \right) \\ &= \text{sum}_{col-wise} \left( \frac{dE(x)}{dy} \right) \end{aligned}$$



Visual Computing  
Institute

RWTH AACHEN  
UNIVERSITY

## Example: Linear/Fully Connected Module - Batching



Given: Derivatives with respect to outputs  $\frac{dE(x^k)}{dy^k}$

Calculate:

Plural!

► Derivatives with respect to inputs

$$\begin{aligned} \frac{dE(x)}{dx} &= \left[ \frac{dE(x^k)}{dx^k} \right]_k = \left[ \frac{dE(x^k)}{dy^k} \cdot \frac{dy^k}{dx^k} \right]_k \\ &= \left[ \frac{dE(x^k)}{dy^k} \cdot W^T \right]_k = \frac{dE(x)}{dy} \cdot W^T \end{aligned}$$

Deriv. w.r.t outputs  
assumed to be  
given row-wise:

$$\frac{dE(x)}{dy} = \begin{bmatrix} - & \frac{dE(x^1)}{dy^1} & - \\ & \vdots & \\ - & \frac{dE(x^n)}{dy^n} & - \end{bmatrix}$$



Visual Computing  
Institute

**RWTH**AACHEN  
UNIVERSITY

## Example: Training a Network

```
1.  network = [module1, module2, ..., modulen], loss = floss
2.
3.  for X, T in batched(inputs, labels) do
4.      z = X
5.      for module in net do
6.          z = module.fprop(z)
7.      end for
8.      E = loss.fprop(z, T)
9.      dz = loss.bprop(1/batchSize)           // Normalization for batch size
10.     for module in reversed(net) do
11.         dz = module.bprop(dz)
12.     end for
13.     for module in net do
14.         module.update(rate)
15.     end for
16. end for
```



# Debugging Tip: Gradient Checking

Check the Jacobian  $J$  from *bprop* with numerical differentiation

- ▶ Numerical approach: Column-wise (here for the first column)

$$\begin{aligned} x_+ &= (x_1 + h_1, x_2, \dots, x_n) \\ x_- &= (x_1 - h_1, x_2, \dots, x_n) \end{aligned} \quad J_{-,1} = \frac{fprop(x_+) - fprop(x_-)}{2h_1}$$

- ▶ Backprop: Row-wise (here for the first row)

$$\begin{aligned} &fprop(x) \\ J_{1,-} &= bprop(1, 0, \dots, 0) \end{aligned}$$

- ▶ Advice

- ▶ Use (small) random  $x$

$$h_i = \sqrt{\epsilon} * \max(x_i, 1)$$



# Expected Results/Tips for MNIST

- ▶ [Linear(28x28, 10), Softmax]
  - ▶ should give  $\pm 750$  errors
- ▶ [Linear(28x28, 200), tanh, Linear(200,10), Softmax]
  - ▶ should give  $\pm 250$  errors
- ▶ Typical learning rates
  - ▶  $\lambda \in [0.1, 0.01]$
- ▶ Typical batch sizes
  - ▶  $N_B \in [100, 1000]$
- ▶ Weight initialization
  - ▶  $W \in \mathbb{R}^{M \times N}$
  - ▶  $W \sim N(0, \sqrt{\frac{2}{M+N}})$ , i.e., sampled from normal distribution around 0 with deviation  $\sqrt{\frac{2}{M+N}}$
  - ▶  $b = 0$
- ▶ Pre-process the data
  - ▶ Divide values by 255 (= max pixel value)

