

Exercise 1: Solutions

Lecture, Summer 2019

Dr. Christian Terboven
Simon Schwitanski

Task 1

- Assume a parallel program with a serial runtime $t = 100\text{s}$ and a parallel fraction $p = 0.95$. What is the minimum runtime and maximum speedup that can be achieved with this program?
- **Serial fraction: $s = 1 - p = 0.05$**
 - Fraction which cannot be parallelized
- **Minimum runtime: 5s**
 - Parallelized program's execution time: $T(N) = (s + \frac{p}{N}) \cdot T(1)$
- **Maximum speedup: $100\text{s} / 5\text{s} = 20$**
- **Speedup according to Amdahl's Law:**

$$S_{0.95}(N) = \frac{T(1)}{T(N)} = \frac{1}{s + \frac{1-s}{N}} = \frac{1}{0.05 + \frac{1-0.05}{N}}$$


- Perform a limit value consideration of S_p for $N \rightarrow \infty$, with N being the number of processors used. What does the result imply regarding the efficiency E_p ?

$$S_p(N) = \frac{T(1)}{T(N)} = \frac{1}{s + \frac{1-s}{N}}$$

$$S_{max} = \lim_{N \rightarrow \infty} S_p(N) = \lim_{N \rightarrow \infty} \frac{1}{s + \frac{1-s}{N}} = \frac{1}{s}$$

$$\lim_{N \rightarrow \infty} \varepsilon_p(N) = \lim_{N \rightarrow \infty} \frac{1}{s(N-1) + 1} = 0$$

Example:
 $s=0.05$



- **Which problems of real-world applications do also limit the achievable speedup, but are not taken into account by Amdahl's Law?**

- **Selection of possible answers:**
 - Overhead of parallelization
 - Creation of threads, Communication of information, ...
 - Contention on shared resources
 - Overhead of locks, ...
 - Additional work with increasing number of processing elements
 - Increasing effort of problem partitioning / load balancing, ...

Task 2

■ Completion of the program skeleton

```
1  #include <iostream>
2  #include <thread>
3
4  void hello()
5  {
6      std::cout << " Hello " << std::endl;
7  }
8
9  void world()
10 {
11     std::cout << " World. " << std::endl;
12 }
```

```
1  int main()  
2  {  
3      // call with two new threads  
4      std::thread t1( hello );  
5      std::thread t2( world );  
6  
7      // join the threads with the main thread  
8      t1.join();  
9      t2.join();  
10  
11     return 0;  
12 }
```


Task 3

■ Construct a scenario in which the naive lock impl. deadlocks

Schedule:

```
bool flag[2];    // init. w/ false
void lock() {
    flag[i] = true;
    while (flag[j]) {}
}
```

Thread 0	Thread 1
flag[0] = true	
	flag[1] = true
	while (flag[0]) {}
while (flag[1]) {}	

- Thread 0 first sets `flag[0]` to true and might be interrupted or just slower than Thread 1
- Thread 1 sets `flag[1]` to true (faster than Thread 0) and observes `flag[0] == true`, so it has to wait
- When thread 0 continues, it observes `flag[1] == true`, so it has to wait
- Both threads wait for each other: Deadlock

■ Does the other naive lock impl. provide mutual exclusion?

- In case of 2 threads: Yes. The variable `victim` is either 0* or 1* and it can never be the case that both threads pass the while condition.
- In case of 3 or more threads: No. Only one thread can be the `victim`, therefore 2 or more threads can simultaneously enter the critical region.
- ***But:** We have a data race on the variable `victim`. According to the C++ standard, this is *undefined behavior*. That means: *Anything* can happen during runtime (e.g., `victim` could be set to a random value or the program could crash). Thus, according to the C++ standard, the implementation obviously does not implement mutual exclusion.
 - Solution to avoid undefined behavior: Make accesses on `victim` atomic.

■ Under which condition does this implementation stop making progress?

→ Single-threaded execution / the other thread terminated

■ Why is this condition not a deadlock?

→ A deadlock requires at least two processes / threads waiting for each other.

Here, we have only one thread waiting for an event.

■ Data Race in `access_one()` ?

→ No.

■ Data Race in `access_two()` ?

→ Yes, two threads potentially write to same `A[i]` simultaneously (at least 11 threads needed here)

■ Data Race in `access_three()` ?

→ Yes, two threads potentially write to same `A[i]` simultaneously (at least 2 threads needed here).

Task 4

■ Why does the `empty()` function have to acquire a lock?

- If not, the access to member variable `data_queue` is not safe.
- Problem: The implementation of `std::queue::empty()` is not guaranteed to be thread-safe (although it might be in some C++ STL implementations).

- **Note:** The `conditional_variable::wait()` function cannot be used with a `lock_guard`, but instead `unique_lock` has to be used. The reason is that the `wait()` function has to `unlock()` and `lock()` a given mutex which is not supported by `lock_guard`.

■ Implementation of the `try_pop()` member function

```
1  bool threadsafe_queue::try_pop(T& value)
2  {
3      std::lock_guard<std::mutex> lk(mut);
4      if (data_queue.empty())
5          return false;
6      value = data_queue.front();
7      data_queue.pop();
8      return true;
9  }
```

- **What happens if some thread calls `push()` with `cond.notify_one()` to notify a waiting thread `j`? Is it possible that another thread `k` gets the released mutex of thread `i` *before* thread `j` which has been woken up?**

→ Yes (see C++ standard on condition variables). But this is not a problem:

Even if thread `k` modifies the queue, thread `j` can only continue execution after it has acquired the mutex. Therefore, it has to wait for thread `k` to release the lock. As soon as thread `j` can go on with execution, it first checks the predicate `(!queue.empty())`. If the queue is empty again, then thread `j` goes again to sleep and waits for a notification. Otherwise, thread `j` can get the element out of the queue.

→ Thus, the implementation is thread-safe.